

myfind

Bo Xuan Yang, Barnabás Körömi

Overview

This C++ program searches for files in a directory and its subdirectories (if enabled) using parallel processes. Each filename is searched in a separate child process, utilizing `fork()`. The program can perform case-insensitive searches and handle recursion with `-R` and `-i` flags. The parent process waits for all child processes to finish before terminating.

Key Parallelization Concepts

Forking Processes:

The program creates a new child process for each filename search using `fork()`. This allows multiple filenames to be searched in parallel.

Process Synchronization:

The parent process waits for all child processes to complete using `wait()`, preventing orphaned "zombie" processes.

Output Handling:

Each child process prints its search results to `stdout`. Since multiple processes output concurrently, the results may be interleaved, but each result includes the PID, making it identifiable.

Functions Explanation

`toLowerCase(std::string &str):`

Converts the input string to lowercase for case-insensitive comparison using `std::transform()`.

`findFile(const std::string &directory, const std::string &filename, bool recursive, bool caseInsensitive):`

Searches for a filename in a specified directory. It handles recursion, case-insensitive comparison, and prints matching files' paths along with the process ID.

`main(int argc, char *argv[]):`

Handles command-line arguments (-R for recursion, -i for case-insensitivity), then forks a child process for each filename. The parent process waits for all children to finish.

Execution Flow

Input Parsing:

`getopt()` processes flags (-R, -i) and gathers the directory and filenames.

Forking:

The program forks a child process for each filename to search in parallel.

File Search:

Each child process calls `findFile()` to search the directory. It handles recursion and case-insensitivity as needed.

Synchronization:

The parent waits for all child processes to complete using `wait()`.

Output:

Each child process outputs its results with the format: `<pid>: <filename>: <full path>`.

Challenges

Concurrency in Output:

Since multiple processes write to `stdout` simultaneously, the output can be interleaved, but each result includes a PID to identify the source process.

Fork Overhead:

Forking a large number of processes may incur overhead, especially if the directory structure is not large.

Build and Run:

Build it within the WSL

```
g++ -o myfind myfind.cpp
```

Test cases

Test Case 1 (pass): Search for a specific file in a directory (non-recursive, case-sensitive):

```
./myfind ./test_dir Test.txt test.doc test
```

Test Case 1 (fail – no output): Search for a specific file in a directory (non-recursive, case-sensitive):

```
./myfind ./test_dir test.txt test.doc test
```

Test Case 2: Search for a file recursively in subdirectories (-R flag):

```
./myfind -R ./test_dir testing.txt
```

Test Case 3: Search for a file with case-insensitive comparison (-i flag):

```
./myfind -i ./test_dir test.txt test.doc test
```

Test Case 4: Combine both -R and -i flags for a recursive, case-insensitive search:

```
./myfind -Ri ./test_dir test.txt test.doc Testing
```